

Project Proposal By Durga Sandeep

Custom Game Environment for Reinforcement Learning

This project aims to build an intelligent agent that learns to navigate and perform tasks in a custom-designed game environment, such as a GridWorld or Maze. The agent interacts with the environment by taking actions (e.g., move up, down, left, right) and receives rewards based on its performance, such as reaching a goal, collecting items, or avoiding obstacles.

Using Reinforcement Learning (RL) techniques like Q-Learning or SARSA, the agent learns an optimal policy to maximize cumulative rewards. The project demonstrates key RL concepts including states, actions, rewards, policies, and exploration vs. exploitation.

1: Imports and Setup

Run this cell first to import all required libraries



```
In [1]: import pygame
import random
import numpy as np
import matplotlib.pyplot as plt
```

```
print("All libraries imported successfully!")
```

pygame 2.6.1 (SDL 2.28.4, Python 3.13.3)

Hello from the pygame community. <https://www.pygame.org/contribute.html>

All libraries imported successfully!

2: Define GridWorld Environment

```
In [2]: class GridWorld:
        """
        A grid-based environment where an agent navigates from start (0,0) to goal (
        Obstacles are randomly placed and block the agent's movement.
        """

        def __init__(self, grid_size=20, cell_size=30):
            """
            Initialize the GridWorld environment.

            Args:
                grid_size: Number of cells in each dimension (creates grid_size x gr
                cell_size: Size of each cell in pixels for rendering
            """
            self.grid_size = grid_size
            self.cell_size = cell_size
            # Set obstacles to 15% of total grid cells
            self.num_obstacles = int(grid_size * grid_size * 0.15)
            self._generate_world()
            self.reset()

        def _generate_world(self):
            """
            Generate random obstacle positions and set goal position.
            Ensures obstacles don't block start (0,0) or goal positions.
            """
            self.obstacles = []
            while len(self.obstacles) < self.num_obstacles:
                pos = [random.randint(0, self.grid_size-1),
                      random.randint(0, self.grid_size-1)]
                # Ensure obstacle is not at start or goal position
                if pos not in self.obstacles and pos != [0,0] and pos != [self.grid_
                    self.obstacles.append(pos)
            # Goal is always at bottom-right corner
            self.goal_pos = [self.grid_size-1, self.grid_size-1]

        def reset(self):
            """
            Reset agent to starting position (0,0).

            Returns:
                tuple: Starting position as (x, y)
            """
            self.agent_pos = [0, 0]
            return tuple(self.agent_pos)
```

```

def step(self, action):
    """
    Execute one action in the environment.

    Args:
        action: String indicating direction ('up', 'down', 'left', 'right')

    Returns:
        tuple: (next_state, reward, done)
            - next_state: New position as (x, y)
            - reward: Numerical reward for the action
            - done: Boolean indicating if goal was reached
    """
    x, y = self.agent_pos
    next_x, next_y = x, y

    # Calculate next position based on action
    if action == 'up':
        next_y = max(y-1, 0)
    elif action == 'down':
        next_y = min(y+1, self.grid_size-1)
    elif action == 'left':
        next_x = max(x-1, 0)
    elif action == 'right':
        next_x = min(x+1, self.grid_size-1)

    done = False

    # Check if agent hit an obstacle
    if [next_x, next_y] in self.obstacles:
        # Large negative reward for hitting obstacle
        reward = -10
        # Agent stays in current position (doesn't move into obstacle)
    else:
        # Valid move - update agent position
        self.agent_pos = [next_x, next_y]

        # Check if agent reached the goal
        if self.agent_pos == self.goal_pos:
            # Large positive reward for reaching goal
            reward = 100
            done = True
        else:
            # Small penalty for each step to encourage shorter paths
            reward = -1

    return tuple(self.agent_pos), reward, done

def render(self, screen, path=None, agent_pos=None):
    """
    Render the grid world using Pygame.

    Args:
        screen: Pygame display surface
        path: Optional list of positions to highlight as the path
        agent_pos: Optional position to render the agent
    """
    for y in range(self.grid_size):
        for x in range(self.grid_size):
            # Create rectangle for this cell

```

```

rect = pygame.Rect(x*self.cell_size, y*self.cell_size,
                   self.cell_size, self.cell_size)

# Draw grid background (dark gray)
pygame.draw.rect(screen, (30, 30, 30), rect)
pygame.draw.rect(screen, (50, 50, 50), rect, 1) # Grid Lines

# Draw obstacles in blue
if [x, y] in self.obstacles:
    pygame.draw.rect(screen, (0, 0, 255), rect)

# Draw goal in green
if [x, y] == self.goal_pos:
    pygame.draw.rect(screen, (0, 255, 0), rect)

# Draw path in yellow
if path and (x, y) in path:
    pygame.draw.rect(screen, (255, 255, 0), rect)

# Draw agent in red (properly handles tuple comparison)
if agent_pos and list(agent_pos) == [x, y]:
    pygame.draw.rect(screen, (255, 0, 0), rect)

```

3: Define Q-Learning Agent

```

In [3]: class QLearningAgent:
        """
        Q-Learning agent that learns optimal policy through exploration and exploitation
        Uses epsilon-greedy strategy for action selection.
        """

        def __init__(self, grid_size=20, actions=['up', 'down', 'left', 'right'],
                      alpha=0.1, gamma=0.95, epsilon=1.0):
            """
            Initialize Q-Learning agent.

            Args:
                grid_size: Size of the grid world
                actions: List of possible actions
                alpha: Learning rate (how much new info overrides old)
                gamma: Discount factor (importance of future rewards)
                epsilon: Initial exploration rate (probability of random action)
            """
            self.grid_size = grid_size
            self.actions = actions
            self.alpha = alpha # Learning rate
            self.gamma = gamma # Discount factor
            self.epsilon = epsilon # Exploration rate
            self.epsilon_min = 0.01 # Minimum exploration rate
            self.epsilon_decay = 0.995 # Rate at which epsilon decreases

            # Q-table: stores expected reward for each (state, action) pair
            # Shape: (grid_size, grid_size, num_actions)
            self.Q = np.zeros((grid_size, grid_size, len(actions)))

        def choose_action(self, state):

```

```

"""
Choose action using epsilon-greedy policy.

Args:
    state: Current state as (x, y) tuple

Returns:
    str: Selected action
"""
x, y = state

# Exploration: choose random action with probability epsilon
if random.random() < self.epsilon:
    return random.choice(self.actions)
else:
    # Exploitation: choose action with highest Q-value
    max_q = np.max(self.Q[x, y])
    # Handle ties by randomly selecting among best actions
    max_actions = [i for i, q in enumerate(self.Q[x, y]) if q == max_q]
    return self.actions[random.choice(max_actions)]

def learn(self, state, action, reward, next_state):
    """
    Update Q-value using the Q-learning update rule.
     $Q(s,a) \leftarrow Q(s,a) + \alpha[r + \gamma \max_{a'} Q(s',a')] - Q(s,a)$ 

    Args:
        state: Current state (x, y)
        action: Action taken
        reward: Reward received
        next_state: Resulting state (x', y')
    """
    x, y = state
    nx, ny = next_state
    idx = self.actions.index(action)

    # Calculate TD target: reward + discounted max future Q-value
    target = reward + self.gamma * np.max(self.Q[nx, ny])

    # Update Q-value using learning rate
    self.Q[x, y, idx] += self.alpha * (target - self.Q[x, y, idx])

```

4: Path Generation Function

```

In [4]: def generate_learned_path(env, agent):
    """
    Generate a path from start to goal using the learned Q-table.
    Uses greedy policy (always choose best action) without exploration.

    Args:
        env: GridWorld environment
        agent: Trained Q-Learning agent

    Returns:
        list: Path as list of (x, y) positions from start to goal
    """

```

```

state = env.reset()
path = [state]
visited = set()
visited.add(state)
max_path_length = 1000 # Prevent infinite loops

while state != tuple(env.goal_pos) and len(path) < max_path_length:
    x, y = state

    # Get all Q-values for current state
    q_values = agent.Q[x, y].copy()

    # Try actions in order of Q-value (best first)
    action_indices = np.argsort(q_values)[::-1] # Sort descending
    moved = False

    for idx in action_indices:
        action = agent.actions[idx]
        nx, ny = x, y

        # Calculate next position for this action
        if action == 'up':
            ny = max(y-1, 0)
        elif action == 'down':
            ny = min(y+1, env.grid_size-1)
        elif action == 'left':
            nx = max(x-1, 0)
        elif action == 'right':
            nx = min(x+1, env.grid_size-1)

        # Skip if next position is an obstacle
        if [nx, ny] in env.obstacles:
            continue

        # Skip if no actual movement occurred
        if (nx, ny) == state:
            continue

        next_state = (nx, ny)

        # Prefer unvisited states, but allow revisiting if necessary
        if next_state not in visited or not moved:
            path.append(next_state)
            if next_state not in visited:
                visited.add(next_state)
            state = next_state
            moved = True
            break

    if not moved:
        # Agent is stuck - cannot find valid move
        print("Warning: Agent got stuck before reaching goal")
        break

# Report path generation results
if state == tuple(env.goal_pos):
    print(f"Success! Path reached goal in {len(path)} steps")
else:
    print(f"Warning: Path did not reach goal. Ended at {state}, goal is {tuple(env.goal_pos)}")

```

```
return path
```

5: Learning Curve Plotting Function

```
In [5]: def plot_learning_curve(episode_rewards, episode_lengths, success_history):
        """
        Create and save comprehensive learning curve visualization.

        Args:
            episode_rewards: List of total rewards per episode
            episode_lengths: List of steps taken per episode
            success_history: List of success rates over time
        """
        fig, axes = plt.subplots(2, 2, figsize=(12, 10))

        # ===== Plot 1: Episode Rewards =====
        axes[0, 0].plot(episode_rewards, alpha=0.6, color='blue')

        # Add moving average to smooth the curve
        window = 50
        if len(episode_rewards) >= window:
            moving_avg = np.convolve(episode_rewards, np.ones(window)/window, mode='
            axes[0, 0].plot(range(window-1, len(episode_rewards)), moving_avg,
                            color='red', linewidth=2, label=f'{window}-episode MA')
            axes[0, 0].legend()
        axes[0, 0].set_xlabel('Episode')
        axes[0, 0].set_ylabel('Total Reward')
        axes[0, 0].set_title('Episode Rewards Over Time')
        axes[0, 0].grid(True, alpha=0.3)

        # ===== Plot 2: Episode Lengths =====
        axes[0, 1].plot(episode_lengths, alpha=0.6, color='green')

        # Add moving average
        if len(episode_lengths) >= window:
            moving_avg = np.convolve(episode_lengths, np.ones(window)/window, mode='
            axes[0, 1].plot(range(window-1, len(episode_lengths)), moving_avg,
                            color='red', linewidth=2, label=f'{window}-episode MA')
            axes[0, 1].legend()
        axes[0, 1].set_xlabel('Episode')
        axes[0, 1].set_ylabel('Steps to Goal')
        axes[0, 1].set_title('Episode Length Over Time')
        axes[0, 1].grid(True, alpha=0.3)

        # ===== Plot 3: Success Rate =====
        axes[1, 0].plot(success_history, color='purple', linewidth=2)
        axes[1, 0].set_xlabel('Episode')
        axes[1, 0].set_ylabel('Success Rate (%)')
        axes[1, 0].set_title(f'Success Rate (Window: {window} episodes)')
        axes[1, 0].grid(True, alpha=0.3)
        axes[1, 0].set_ylim([0, 105])

        # ===== Plot 4: Summary Statistics =====
        axes[1, 1].axis('off')
        summary_text = f"""
```

Training Summary

Total Episodes: {len(episode_rewards)}

Final Success Rate: {success_history[-1]:.1f}%

Average Reward (last 100):

{np.mean(episode_rewards[-100:]):.2f}

Average Length (last 100):

{np.mean(episode_lengths[-100:]):.2f}

Best Episode Reward: {max(episode_rewards):.2f}

Worst Episode Reward: {min(episode_rewards):.2f}

"""

axes[1, 1].text(0.1, 0.5, summary_text, fontsize=12, family='monospace',
verticalalignment='center')

plt.tight_layout()

plt.savefig('learning_curve.png', dpi=150, bbox_inches='tight')

print("\nLearning curve saved as 'learning_curve.png'")

plt.close()

6: Main Training and Visualization

```
In [6]: def main():
        """
        Main function to train Q-Learning agent and visualize results.
        """

        # ===== Environment Configuration =====
        GRID_SIZE = 20
        CELL_SIZE = 30
        WINDOW_SIZE = GRID_SIZE * CELL_SIZE

        # ===== Initialize Pygame =====
        pygame.init()
        screen = pygame.display.set_mode((WINDOW_SIZE, WINDOW_SIZE))
        pygame.display.set_caption("GridWorld Q-Learning Animated Path")
        clock = pygame.time.Clock()

        # ===== Create Environment and Agent =====
        env = GridWorld(GRID_SIZE, CELL_SIZE)
        agent = QLearningAgent(GRID_SIZE)

        print("Training agent...")

        # ===== Training Configuration =====
        num_episodes = 2000      # Total number of training episodes
        max_steps = 500         # Maximum steps per episode
        success_count = 0       # Counter for successful episodes

        # Lists to track training progress
        episode_rewards = []
        episode_lengths = []
        success_history = []
```



```

# Track state visitation frequency for heatmap
visit_map = np.zeros((GRID_SIZE, GRID_SIZE))

# ===== TRAINING LOOP =====
for episode in range(num_episodes):
    state = env.reset()
    done = False
    step = 0
    total_reward = 0

    # Run episode until goal reached or max steps
    while not done and step < max_steps:
        step += 1

        # Record state visit for heatmap
        visit_map[state[0], state[1]] += 1

        # Agent chooses and executes action
        action = agent.choose_action(state)
        next_state, reward, done = env.step(action)

        # Agent learns from experience
        agent.learn(state, action, reward, next_state)

        total_reward += reward
        state = next_state

    # Record episode statistics
    episode_rewards.append(total_reward)
    episode_lengths.append(step)

    if done:
        success_count += 1

    # Calculate rolling success rate
    window = 50
    if episode >= window:
        recent_success = sum(1 for i in range(episode - window + 1, episode)
                             if episode_lengths[i] < max_steps and episode_rewards[i] > 0)
        success_rate = (recent_success / window) * 100
    else:
        success_rate = (success_count / (episode + 1)) * 100
    success_history.append(success_rate)

    # Decay exploration rate
    if agent.epsilon > agent.epsilon_min:
        agent.epsilon *= agent.epsilon_decay

    # Progress update every 200 episodes
    if (episode + 1) % 200 == 0:
        print(f"Episode {episode + 1}/{num_episodes}, Success Rate: {success_rate}%")

print(f"\nTraining complete! Successes: {success_count}/{num_episodes}")
print(f"Final epsilon: {agent.epsilon:.3f}")

# ===== Save Learning Curve =====
plot_learning_curve(episode_rewards, episode_lengths, success_history)

# ===== Generate and Save Heatmap =====
plt.imshow(visit_map, cmap='hot', interpolation='nearest')

```

```

plt.colorbar(label="Visit Frequency")
plt.title("Q-Learning State Visit Heatmap")
plt.savefig("visit_heatmap.png", dpi=200)
plt.close()
print("Heat map saved as visit_heatmap.png")

# ===== Generate Learned Path =====
print("Generating learned path...")
path = generate_learned_path(env, agent)
print(f"Path length = {len(path)}")

# ===== ANIMATE PATH VISUALIZATION =====
print("\nAnimating final path frame-by-frame...\n")

# Draw path step by step
for i, pos in enumerate(path):
    screen.fill((0, 0, 0))
    # Render grid with path up to current position
    env.render(screen, path=path[:i+1], agent_pos=pos)
    pygame.display.flip()
    pygame.time.delay(120) # Delay in ms (smaller = faster animation)

# ===== Hold Display Until User Closes Window =====
running = True
while running:
    for event in pygame.event.get():
        if event.type == pygame.QUIT:
            running = False
    clock.tick(30)

pygame.quit()

```

7: Run the Program

```

In [7]: if __name__ == "__main__":
        main()

```

```
Training agent...
Episode 200/2000, Success Rate: 0.0%, Epsilon: 0.367
Episode 400/2000, Success Rate: 26.0%, Epsilon: 0.135
Episode 600/2000, Success Rate: 82.0%, Epsilon: 0.049
Episode 800/2000, Success Rate: 100.0%, Epsilon: 0.018
Episode 1000/2000, Success Rate: 100.0%, Epsilon: 0.010
Episode 1200/2000, Success Rate: 100.0%, Epsilon: 0.010
Episode 1400/2000, Success Rate: 100.0%, Epsilon: 0.010
Episode 1600/2000, Success Rate: 100.0%, Epsilon: 0.010
Episode 1800/2000, Success Rate: 100.0%, Epsilon: 0.010
Episode 2000/2000, Success Rate: 100.0%, Epsilon: 0.010
```

```
Training complete! Successes: 1917/2000
Final epsilon: 0.010
```

```
Learning curve saved as 'learning_curve.png'
Heat map saved as visit_heatmap.png
Generating learned path...
Success! Path reached goal in 39 steps
Path length = 39
```

```
Animating final path frame-by-frame...
```

8: Summary

This notebook demonstrates:

1. **GridWorld Environment:** A customizable grid with obstacles and a goal
2. **Q-Learning Agent:** Learns optimal navigation policy through trial and error
3. **Training Visualization:** Plots showing learning progress over time
4. **Heatmap:** Shows which states were visited most during training
5. **Path Animation:** Visualizes the learned path from start to goal

Key Hyperparameters:

- Learning rate (α): 0.1
- Discount factor (γ): 0.95
- Initial exploration rate (ϵ): 1.0
- Epsilon decay: 0.995


Color Scheme:

-  Red = Agent
-  Green = Goal
-  Blue = Obstacles
-  Yellow = Path taken

9: HEATMAP IMAGE


```
In [8]: import matplotlib.pyplot as plt
import matplotlib.image as mpimg

img = mpimg.imread("visit_heatmap.png")
plt.imshow(img)
plt.axis('off') # Hide axis for clean output
plt.show()
```



10: LEARNING CURVE IMAGE

```
In [9]: img = mpimg.imread("learning_curve.png")
plt.imshow(img)
plt.axis('off') # Hide axis for clean output
plt.show()
```



11: It fails when we go for a huge grid size

```
In [10]: from IPython.display import Video
Video("VideoRec1.mp4", width=600)
```

Out[10]:



12: What happened ?

Q-Learning for RL Agents with Dynamic Obstacles

Can Q-learning handle dynamic obstacles?

Yes, but with important limitations based on environment complexity.

When Q-Learning Works

- **Small, discrete state spaces** (e.g., gridworld)
 - **Obstacle positions** can be encoded in the state
 - **Low-dimensional** environments
-

When Q-Learning Struggles

- **Large state spaces** (images, lidar, continuous positions)
- **Highly dynamic** or continuous environments
- **Many obstacles** or moving agents

Why? Q-tables grow exponentially with state/action dimensions.

Making Q-Learning Work

1. Gridworld Approach

- **State**: agent position + obstacle positions
- **Q-table**: stores $Q(s, a)$ for all combinations
- **Feasible**: for small grids (e.g., 5×5)

2. State Abstraction

Reduce state space by encoding only:

- Distance to nearest obstacle
- Relative direction
- Safety of next cell

3. Reward Shaping

```
rewards = {  
    'collision': -100,  
    'timestep': -1,  
    'goal': +100,  
    'avoid_close_call': +5  
}
```

Algorithm Comparison

Environment Type	Q-learning	DQN	PPO/SAC
Small gridworld	Good	Optional	Not needed
Dynamic gridworld	Limited	Better	Not needed
Continuous	No	Maybe	Best
Images/sensors	No	Good	Good

Key Takeaway

Q-learning works for small, discrete environments with dynamic obstacles.

For complex scenarios → use **DQN** or **PPO/SAC**.

13: Next

Reinforcement Learning Algorithms: DQN, PPO, and SAC

A comprehensive comparison of three popular deep RL algorithms.

1. DQN (Deep Q-Network)

Type: Value-based, off-policy RL

Action Space: Discrete only

Core Idea: Learn a Q-function $Q(s, a)$ that estimates the expected return for taking action a in state s .

Key Features

- Uses a neural network to approximate the Q-table
- **Off-policy:** Learns from past experiences stored in a replay buffer
- **Target network:** Stabilizes learning with a slowly updated copy of the Q-network
- Works well for discrete action spaces (e.g., Atari games)

Limitations

- Doesn't naturally work with continuous actions
- Can be unstable for very large or complex state spaces
- Exploration typically handled via ϵ -greedy strategy

2. PPO (Proximal Policy Optimization)

Type: Policy-based, on-policy RL

Action Space: Discrete or continuous

Core Idea: Directly learn a policy $\pi_{\theta}(a|s)$ to maximize expected reward while preventing large destabilizing updates.

Key Features

- Uses a **clipped surrogate objective** to keep policy updates small
- **On-policy:** Learns from the current policy's experience
- More stable than vanilla policy gradient methods
- Handles both continuous and discrete actions

Strengths

- Good balance between performance and simplicity
- Robust for complex environments
- Widely used in robotics and game AI

Limitations

- Less sample-efficient than off-policy methods
 - Requires more environment interactions
-

3. SAC (Soft Actor-Critic)

Type: Actor-Critic, off-policy, entropy-regularized RL

Action Space: Continuous

Core Idea: Learn both a policy (actor) and Q-function (critic) while encouraging exploration via an entropy term.

Key Features

- Maximizes **reward + entropy** to encourage exploration
- **Off-policy:** Uses replay buffer for sample efficiency
- Supports continuous action spaces naturally
- Uses **two Q-networks** (Double Q-Learning) to reduce overestimation bias

Strengths

- Very stable and sample-efficient for continuous control
- Excellent for robotics, control tasks (MuJoCo, PyBullet)

Limitations

- More complex to implement than PPO
- Slower per update due to multiple networks

Comparison Table

Algorithm	Type	Policy/Value	Action Space	On/Off-policy	Key Use Cases
DQN	Value-based	Q-function	Discrete	Off-policy	Atari games, grid-world
PPO	Policy-based	Policy	Discrete/Continuous	On-policy	Robotics, game AI
SAC	Actor-Critic	Actor + Critic	Continuous	Off-policy	Continuous control, robotics

Key Takeaways

- **DQN**: Simple Q-learning approach for discrete actions, uses experience replay for stability
- **PPO**: Stable policy optimization that works for both action types, preferred for its simplicity and robustness
- **SAC**: Sample-efficient algorithm for continuous control with built-in exploration via entropy maximization

When to Use What?

- **Discrete actions (games, navigation)** → DQN or PPO
 - **Continuous control (robotics, physics)** → PPO or SAC
 - **Need sample efficiency** → DQN or SAC (off-policy)
 - **Need simplicity and stability** → PPO
-

The links that I referred for the approaches which fit better to grid generation with dynamic obstacles - mentioning few of the links below

[Click here to open reference link 1 from Playing Atari with Deep Reinforcement Learning](#)

[Click here to open reference link 2 from www.reinforcementlearningpath.com](#)

[Click here to open reference link 3 from www.findingtheta.com](#)

In []: